

04b_Modeling_Final

April 14, 2023

1 ADS-508-01-SP23 Team 8: Final Project

2 Train model

Much of the code is modified from Fregly, C., & Barth, A. (2021). Data science on AWS: Implementing end-to-end, continuous AI and machine learning pipelines. O'Reilly.

2.1 Install missing dependencies

[PyAthena](#) is a Python DB API 2.0 (PEP 249) compliant client for Amazon Athena.

```
[2]: !pip install --disable-pip-version-check -q PyAthena==2.1.0
      !pip install --disable-pip-version-check -q sagemaker-experiments==0.1.26
      !pip install missingno
      !pip install scikit-optimize
```

WARNING: The directory '/root/.cache/pip' or its parent directory is not owned or is not writable by the current user. The cache has been disabled. Check the permissions and owner of that directory. If executing pip with sudo, you should use sudo's -H flag.

WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead:

<https://pip.pypa.io/warnings/venv>

WARNING: The directory '/root/.cache/pip' or its parent directory is not owned or is not writable by the current user. The cache has been disabled. Check the permissions and owner of that directory. If executing pip with sudo, you should use sudo's -H flag.

WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead:

<https://pip.pypa.io/warnings/venv>

WARNING: The directory '/root/.cache/pip' or its parent directory is not owned or is not writable by the current user. The cache has been disabled. Check the permissions and owner of that directory. If executing pip with sudo, you should use sudo's -H flag.

Collecting missingno

Downloading missingno-0.5.2-py3-none-any.whl (8.7 kB)

Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages (from missingno) (1.21.6)

Requirement already satisfied: seaborn in /opt/conda/lib/python3.7/site-packages (from missingno) (0.10.0)

Requirement already satisfied: matplotlib in /opt/conda/lib/python3.7/site-packages (from missingno) (3.1.3)

Requirement already satisfied: scipy in /opt/conda/lib/python3.7/site-packages (from missingno) (1.4.1)

Requirement already satisfied: kiwisolver>=1.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib->missingno) (1.1.0)

Requirement already satisfied: python-dateutil>=2.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib->missingno) (2.8.2)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib->missingno) (2.4.6)

Requirement already satisfied: cycycler>=0.10 in /opt/conda/lib/python3.7/site-

```

packages (from matplotlib->missingno) (0.10.0)
Requirement already satisfied: pandas>=0.22.0 in /opt/conda/lib/python3.7/site-
packages (from seaborn->missingno) (1.3.5)
Requirement already satisfied: six in /opt/conda/lib/python3.7/site-packages
(from cycycler>=0.10->matplotlib->missingno) (1.14.0)
Requirement already satisfied: setuptools in /opt/conda/lib/python3.7/site-
packages (from kiwisolver>=1.0.1->matplotlib->missingno) (59.3.0)
Requirement already satisfied: pytz>=2017.3 in /opt/conda/lib/python3.7/site-
packages (from pandas>=0.22.0->seaborn->missingno) (2019.3)
Installing collected packages: missingno
Successfully installed missingno-0.5.2
WARNING: Running pip as the 'root' user can result in broken permissions
and conflicting behaviour with the system package manager. It is recommended to
use a virtual environment instead: https://pip.pypa.io/warnings/venv
WARNING: The directory '/root/.cache/pip' or its parent directory is
not owned or is not writable by the current user. The cache has been disabled.
Check the permissions and owner of that directory. If executing pip with sudo,
you should use sudo's -H flag.

Collecting scikit-optimize
  Downloading scikit_optimize-0.9.0-py2.py3-none-any.whl (100 kB)
                                100.3/100.3 kB
222.4 MB/s eta 0:00:00
Requirement already satisfied: scikit-learn>=0.20.0 in
/opt/conda/lib/python3.7/site-packages (from scikit-optimize) (0.22.1)
Requirement already satisfied: numpy>=1.13.3 in /opt/conda/lib/python3.7/site-
packages (from scikit-optimize) (1.21.6)
Requirement already satisfied: joblib>=0.11 in /opt/conda/lib/python3.7/site-
packages (from scikit-optimize) (1.2.0)
Requirement already satisfied: scipy>=0.19.1 in /opt/conda/lib/python3.7/site-
packages (from scikit-optimize) (1.4.1)
Collecting pyaml>=16.9
  Downloading pyaml-21.10.1-py2.py3-none-any.whl (24 kB)
Requirement already satisfied: PyYAML in /opt/conda/lib/python3.7/site-packages
(from pyaml>=16.9->scikit-optimize) (6.0)
Installing collected packages: pyaml, scikit-optimize
Successfully installed pyaml-21.10.1 scikit-optimize-0.9.0
WARNING: Running pip as the 'root' user can result in broken permissions
and conflicting behaviour with the system package manager. It is recommended to
use a virtual environment instead: https://pip.pypa.io/warnings/venv

```

2.2 Globally import libraries

```
[3]: import boto3
from botocore.client import ClientError
import pandas as pd
import numpy as np
from pyathena import connect
from IPython.core.display import display, HTML
import missingno as msno
from skopt import BayesSearchCV
from skopt.space import Real, Categorical, Integer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import make_pipeline, Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.impute import SimpleImputer
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.linear_model import Lasso
import datetime as dt
import time
import sagemaker
from smexperiments.experiment import Experiment
from smexperiments.trial import Trial
import joblib
import os
from io import BytesIO

%matplotlib inline
```

2.3 Instantiate AWS SageMaker and S3 sessions

```
[4]: session = boto3.session.Session()
role = sagemaker.get_execution_role()
region = session.region_name
sagemaker_session = sagemaker.Session()
def_bucket = sagemaker_session.default_bucket()
bucket = 'sagemaker-us-east-ads508-sp23-t8'

s3 = boto3.Session().client(service_name="s3",
                             region_name=region)

sm = boto3.Session().client(service_name="sagemaker",
                             region_name=region)
```

```
[5]: print(f"Default bucket: {def_bucket}")
      print(f"Public T8 bucket: {bucket}")
```

Default bucket: sagemaker-us-east-1-657724983756
Public T8 bucket: sagemaker-us-east-ads508-sp23-t8

2.4 Pass in train and test X from CSV

```
[6]: s3_train_x01_csv_path = f"s3://{def_bucket}/team_8_data/modeling_data/training/
      ↪train_x01.csv"
      train_x01 = pd.read_csv(s3_train_x01_csv_path)
      s3_test_x01_csv_path = f"s3://{def_bucket}/team_8_data/modeling_data/testing/
      ↪test_x01.csv"
      test_x01 = pd.read_csv(s3_test_x01_csv_path)

      print(f'{train_x01.shape}')
      print(f'\n{test_x01.shape}')
```

(25284, 48)

(6321, 48)

2.5 Pass in train and test y from np array

```
[7]: # Define the S3 object key
      train_y01_s3_key = 'team_8_data/modeling_data/training/train_y01.npy'

      # Load the numpy array from S3
      with BytesIO() as data:
          s3.download_fileobj(def_bucket, train_y01_s3_key, data)
          data.seek(0)
          train_y01 = np.load(data)

      # Define the S3 object key
      test_y01_s3_key = 'team_8_data/modeling_data/testing/test_y01.npy'

      # Load the numpy array from S3
      with BytesIO() as data:
          s3.download_fileobj(def_bucket, test_y01_s3_key, data)
          data.seek(0)
          test_y01 = np.load(data)

      train_y01 = train_y01.ravel()
      test_y01 = test_y01.ravel()

      # Confirm that the numpy array was loaded from S3
      print(f'{train_y01.shape}')
      print(f'{test_y01.shape}')
```

```
(25284,)
(6321,)
```

2.6 Model Training using Grid search with 5-fold cross-validation

2.6.1 Neural Network

```
[8]: # Start timer script
start_time = dt.datetime.today()

# Citation: Hochberg, 2018; Shanmukh, 2021
m1v1_nn_pip = Pipeline([('si', SimpleImputer(strategy='median')),
                        ('ss', StandardScaler()),
                        ('nn', MLPRegressor(random_state=1699))])

nodes_h = 3
predictors_p = 49

hidden_layer_sizes_hparam = [[100,],
                              [(nodes_h*(predictors_p+1))+nodes_h+1,],
                              [50, 50]
                              ]

activation_hparam = ['logistic', 'relu']
solver_hparam = ['adam']
alpha_hparam = [.0001, .0005, .001]
learn_rate_hparam = ['constant', 'invscaling']

#hidden_layer_sizes_hparam = [[100,]]
#activation_hparam = ['relu']
#solver_hparam = ['adam']
#alpha_hparam = [.0001]
#learn_rate_hparam = ['invscaling']

m1v1_nn_grd = {'nn_hidden_layer_sizes': hidden_layer_sizes_hparam,
               'nn_activation': activation_hparam,
               'nn_solver': solver_hparam,
               'nn_alpha': alpha_hparam,
               'nn_learning_rate': learn_rate_hparam
               }

m1v1_nn = GridSearchCV(m1v1_nn_pip,
                       m1v1_nn_grd,
                       scoring='neg_root_mean_squared_error',
                       n_jobs=2,
                       refit=True,
                       verbose=2)

m1v1_nn.fit(train_x01, train_y01)
```

```

print(f'Best Estimator:\n{m1v1_nn.best_estimator_}')

print(pd.DataFrame(m1v1_nn.cv_results_))

train_m1v1_nn_y01_pred = m1v1_nn.predict(train_x01)
print(train_m1v1_nn_y01_pred)

test_m1v1_nn_y01_pred = m1v1_nn.predict(test_x01)
print(test_m1v1_nn_y01_pred)

# Display evaluation metrics
# R-sq
train_m1v1_nn_r2 = r2_score(train_y01, train_m1v1_nn_y01_pred)
test_m1v1_nn_r2 = r2_score(test_y01, test_m1v1_nn_y01_pred)

print(f'Train R-sq:\n{train_m1v1_nn_r2}')
print(f'Test R-sq:\n{test_m1v1_nn_r2}')

# RMSE
train_m1v1_nn_rmse = mean_squared_error(train_y01, train_m1v1_nn_y01_pred,
    ↪squared=False)
test_m1v1_nn_rmse = mean_squared_error(test_y01, test_m1v1_nn_y01_pred,
    ↪squared=False)

print(f'Train RMSE:\n{train_m1v1_nn_rmse}')
print(f'Test RMSE:\n{test_m1v1_nn_rmse}')

# End timer script
end_time = dt.datetime.today()
time_elapse = end_time - start_time
print(f'End Time = {end_time}')
print(f'Script Time = {time_elapse}')

```

Fitting 5 folds for each of 36 candidates, totalling 180 fits

```

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 37 tasks      | elapsed: 12.0min
[Parallel(n_jobs=2)]: Done 158 tasks    | elapsed: 45.9min
[Parallel(n_jobs=2)]: Done 180 out of 180 | elapsed: 51.7min finished

```

Best Estimator:

```

Pipeline(memory=None,
      steps=[('si',
              SimpleImputer(add_indicator=False, copy=True, fill_value=None,
                             missing_values=nan, strategy='median',
                             verbose=0)),
             ('ss',

```

```

StandardScaler(copy=True, with_mean=True, with_std=True)),
('nn',
MLPRegressor(activation='relu', alpha=0.0005,
              batch_size='auto', beta_1=0.9, beta_2=0.999,
              early_stopping=False, epsilon=1e-08,
              hidden_layer_sizes=[50, 50],
              learning_rate='constant',
              learning_rate_init=0.001, max_fun=15000,
              max_iter=200, momentum=0.9, n_iter_no_change=10,
              nesterovs_momentum=True, power_t=0.5,
              random_state=1699, shuffle=True, solver='adam',
              tol=0.0001, validation_fraction=0.1,
              verbose=False, warm_start=False))],
verbose=False)

```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\
0	33.143548	0.529329	0.032452	0.003388	
1	34.235025	0.891932	0.034432	0.003698	
2	45.957392	0.721042	0.042339	0.000732	
3	46.024776	0.781211	0.050556	0.017036	
4	38.072548	0.445525	0.033994	0.003400	
5	36.995700	0.322889	0.032575	0.000501	
6	32.707927	0.229827	0.032212	0.000665	
7	32.493836	0.077391	0.032501	0.000327	
8	45.085183	0.421935	0.044698	0.003383	
9	45.192096	0.353731	0.045635	0.003949	
10	37.167754	0.503359	0.034419	0.004507	
11	37.187461	0.628572	0.033794	0.003179	
12	32.951811	0.651863	0.032234	0.000757	
13	32.465124	0.270572	0.032322	0.000358	
14	47.043259	2.372003	0.042367	0.000092	
15	45.504390	0.374914	0.042963	0.001024	
16	37.119073	0.681981	0.033486	0.001257	
17	37.297852	0.599174	0.032584	0.000332	
18	25.550653	0.229620	0.020148	0.003253	
19	25.787540	0.218238	0.017209	0.000411	
20	33.771521	0.603930	0.024322	0.009828	
21	33.597681	0.430749	0.019367	0.000481	
22	31.016342	0.396437	0.017053	0.001032	
23	31.029871	0.554800	0.018221	0.001761	
24	26.011305	0.443071	0.017106	0.000348	
25	25.785980	0.471428	0.016555	0.001151	
26	33.899021	0.349528	0.020482	0.001773	
27	33.563256	0.528809	0.019615	0.000963	
28	31.141742	0.566396	0.020337	0.006852	
29	30.997358	0.365988	0.017249	0.000827	
30	25.684840	0.256868	0.017028	0.000479	
31	25.798281	0.302832	0.017244	0.000354	
32	33.996955	0.638432	0.019921	0.000500	

33	33.543362	0.277484	0.019681	0.000308
34	30.799481	0.347098	0.017663	0.000995
35	30.831124	0.491073	0.015790	0.002267

	param_nn_activation	param_nn_alpha	param_nn_hidden_layer_sizes \
0	logistic	0.0001	[100]
1	logistic	0.0001	[100]
2	logistic	0.0001	[154]
3	logistic	0.0001	[154]
4	logistic	0.0001	[50, 50]
5	logistic	0.0001	[50, 50]
6	logistic	0.0005	[100]
7	logistic	0.0005	[100]
8	logistic	0.0005	[154]
9	logistic	0.0005	[154]
10	logistic	0.0005	[50, 50]
11	logistic	0.0005	[50, 50]
12	logistic	0.001	[100]
13	logistic	0.001	[100]
14	logistic	0.001	[154]
15	logistic	0.001	[154]
16	logistic	0.001	[50, 50]
17	logistic	0.001	[50, 50]
18	relu	0.0001	[100]
19	relu	0.0001	[100]
20	relu	0.0001	[154]
21	relu	0.0001	[154]
22	relu	0.0001	[50, 50]
23	relu	0.0001	[50, 50]
24	relu	0.0005	[100]
25	relu	0.0005	[100]
26	relu	0.0005	[154]
27	relu	0.0005	[154]
28	relu	0.0005	[50, 50]
29	relu	0.0005	[50, 50]
30	relu	0.001	[100]
31	relu	0.001	[100]
32	relu	0.001	[154]
33	relu	0.001	[154]
34	relu	0.001	[50, 50]
35	relu	0.001	[50, 50]

	param_nn_learning_rate	param_nn_solver \
0	constant	adam
1	invscaling	adam
2	constant	adam
3	invscaling	adam
4	constant	adam

5	invscaling	adam
6	constant	adam
7	invscaling	adam
8	constant	adam
9	invscaling	adam
10	constant	adam
11	invscaling	adam
12	constant	adam
13	invscaling	adam
14	constant	adam
15	invscaling	adam
16	constant	adam
17	invscaling	adam
18	constant	adam
19	invscaling	adam
20	constant	adam
21	invscaling	adam
22	constant	adam
23	invscaling	adam
24	constant	adam
25	invscaling	adam
26	constant	adam
27	invscaling	adam
28	constant	adam
29	invscaling	adam
30	constant	adam
31	invscaling	adam
32	constant	adam
33	invscaling	adam
34	constant	adam
35	invscaling	adam

	params	split0_test_score \
0	{'nn__activation': 'logistic', 'nn__alpha': 0...	-7.886156
1	{'nn__activation': 'logistic', 'nn__alpha': 0...	-7.886156
2	{'nn__activation': 'logistic', 'nn__alpha': 0...	-6.666789
3	{'nn__activation': 'logistic', 'nn__alpha': 0...	-6.666789
4	{'nn__activation': 'logistic', 'nn__alpha': 0...	-4.942271
5	{'nn__activation': 'logistic', 'nn__alpha': 0...	-4.942271
6	{'nn__activation': 'logistic', 'nn__alpha': 0...	-7.885017
7	{'nn__activation': 'logistic', 'nn__alpha': 0...	-7.885017
8	{'nn__activation': 'logistic', 'nn__alpha': 0...	-6.660854
9	{'nn__activation': 'logistic', 'nn__alpha': 0...	-6.660854
10	{'nn__activation': 'logistic', 'nn__alpha': 0...	-4.943635
11	{'nn__activation': 'logistic', 'nn__alpha': 0...	-4.943635
12	{'nn__activation': 'logistic', 'nn__alpha': 0...	-7.886704
13	{'nn__activation': 'logistic', 'nn__alpha': 0...	-7.886704
14	{'nn__activation': 'logistic', 'nn__alpha': 0...	-6.656225

15	{'nn_activation': 'logistic', 'nn_alpha': 0...	-6.656225
16	{'nn_activation': 'logistic', 'nn_alpha': 0...	-4.945009
17	{'nn_activation': 'logistic', 'nn_alpha': 0...	-4.945009
18	{'nn_activation': 'relu', 'nn_alpha': 0.0001...	-4.755215
19	{'nn_activation': 'relu', 'nn_alpha': 0.0001...	-4.755215
20	{'nn_activation': 'relu', 'nn_alpha': 0.0001...	-3.408820
21	{'nn_activation': 'relu', 'nn_alpha': 0.0001...	-3.408820
22	{'nn_activation': 'relu', 'nn_alpha': 0.0001...	-2.138443
23	{'nn_activation': 'relu', 'nn_alpha': 0.0001...	-2.138443
24	{'nn_activation': 'relu', 'nn_alpha': 0.0005...	-4.793217
25	{'nn_activation': 'relu', 'nn_alpha': 0.0005...	-4.793217
26	{'nn_activation': 'relu', 'nn_alpha': 0.0005...	-3.415002
27	{'nn_activation': 'relu', 'nn_alpha': 0.0005...	-3.415002
28	{'nn_activation': 'relu', 'nn_alpha': 0.0005...	-2.072638
29	{'nn_activation': 'relu', 'nn_alpha': 0.0005...	-2.072638
30	{'nn_activation': 'relu', 'nn_alpha': 0.001,...	-4.732671
31	{'nn_activation': 'relu', 'nn_alpha': 0.001,...	-4.732671
32	{'nn_activation': 'relu', 'nn_alpha': 0.001,...	-3.442814
33	{'nn_activation': 'relu', 'nn_alpha': 0.001,...	-3.442814
34	{'nn_activation': 'relu', 'nn_alpha': 0.001,...	-2.158172
35	{'nn_activation': 'relu', 'nn_alpha': 0.001,...	-2.158172

	split1_test_score	split2_test_score	split3_test_score	\
0	-8.056679	-7.571552	-7.917553	
1	-8.056679	-7.571552	-7.917553	
2	-6.611240	-6.292724	-6.284235	
3	-6.611240	-6.292724	-6.284235	
4	-4.733368	-4.564217	-4.642696	
5	-4.733368	-4.564217	-4.642696	
6	-8.054838	-7.574916	-7.917823	
7	-8.054838	-7.574916	-7.917823	
8	-6.610766	-6.290245	-6.284635	
9	-6.610766	-6.290245	-6.284635	
10	-4.732347	-4.564699	-4.641053	
11	-4.732347	-4.564699	-4.641053	
12	-8.054588	-7.577553	-7.916608	
13	-8.054588	-7.577553	-7.916608	
14	-6.610904	-6.288636	-6.287919	
15	-6.610904	-6.288636	-6.287919	
16	-4.731240	-4.566586	-4.652802	
17	-4.731240	-4.566586	-4.652802	
18	-4.727183	-4.493818	-4.706454	
19	-4.727183	-4.493818	-4.706454	
20	-3.261693	-3.143096	-3.177231	
21	-3.261693	-3.143096	-3.177231	
22	-2.074308	-1.988741	-2.209510	
23	-2.074308	-1.988741	-2.209510	
24	-4.620354	-4.531985	-4.668786	

25	-4.620354	-4.531985	-4.668786
26	-3.165363	-3.111582	-3.153980
27	-3.165363	-3.111582	-3.153980
28	-1.978748	-1.917381	-2.099757
29	-1.978748	-1.917381	-2.099757
30	-4.652343	-4.511500	-4.645985
31	-4.652343	-4.511500	-4.645985
32	-3.205790	-3.099954	-3.160798
33	-3.205790	-3.099954	-3.160798
34	-1.995288	-1.845477	-2.203662
35	-1.995288	-1.845477	-2.203662

	split4_test_score	mean_test_score	std_test_score	rank_test_score
0	-7.868349	-7.860058	0.158714	31
1	-7.868349	-7.860058	0.158714	31
2	-6.383989	-6.447796	0.160961	29
3	-6.383989	-6.447796	0.160961	29
4	-4.418050	-4.660120	0.174970	19
5	-4.418050	-4.660120	0.174970	19
6	-7.869608	-7.860440	0.157032	33
7	-7.869608	-7.860440	0.157032	33
8	-6.380977	-6.445496	0.159897	27
9	-6.380977	-6.445496	0.159897	27
10	-4.419229	-4.660193	0.174983	21
11	-4.419229	-4.660193	0.174983	21
12	-7.871709	-7.861432	0.156003	35
13	-7.871709	-7.861432	0.156003	35
14	-6.378792	-6.444495	0.158522	25
15	-6.378792	-6.444495	0.158522	25
16	-4.421981	-4.663523	0.174169	23
17	-4.421981	-4.663523	0.174169	23
18	-4.317670	-4.600068	0.168870	17
19	-4.317670	-4.600068	0.168870	17
20	-3.227816	-3.243731	0.092065	11
21	-3.227816	-3.243731	0.092065	11
22	-1.943277	-2.070856	0.096772	5
23	-1.943277	-2.070856	0.096772	5
24	-4.332465	-4.589361	0.153697	15
25	-4.332465	-4.589361	0.153697	15
26	-3.171301	-3.203446	0.107820	7
27	-3.171301	-3.203446	0.107820	7
28	-2.024192	-2.018543	0.065378	1
29	-2.024192	-2.018543	0.065378	1
30	-4.367697	-4.582040	0.128561	13
31	-4.367697	-4.582040	0.128561	13
32	-3.234924	-3.228856	0.116271	9
33	-3.234924	-3.228856	0.116271	9
34	-2.029527	-2.046425	0.126879	3

```

35          -2.029527          -2.046425          0.126879          3
[ 5.77257103  2.71321492 34.23584791 ...  2.24031845 14.15423095
 0.35154327]
[41.00555646 -0.83261161 37.19630089 ... -1.23440346  6.85440592
17.10392427]
Train R-sq:
0.9953169265206908
Test R-sq:
0.9946084935331966
Train RMSE:
1.2934970735367237
Test RMSE:
1.4032555490004501
End Time = 2023-04-13 19:29:37.552207
Script Time = 0:52:22.942444

/opt/conda/lib/python3.7/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:571:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
  % self.max_iter, ConvergenceWarning)

```

```

[9]: s3_m1v1_nn_pqt_base_path = f"../models"

if not os.path.exists(s3_m1v1_nn_pqt_base_path):
    os.makedirs(s3_m1v1_nn_pqt_base_path)

s3_m1v1_nn_pqt_path = os.path.join(s3_m1v1_nn_pqt_base_path,
                                   'm1v1_nn.parquet')

# save the model to disk using joblib
joblib.dump(m1v1_nn,
            s3_m1v1_nn_pqt_path)

# load the saved model from disk using joblib
m1v1_nn_fitted = joblib.load(s3_m1v1_nn_pqt_path)

```

```

[10]: # specify the S3 bucket and key where you want to save the model
m1v1_nn_key_name = 'team_8_data/models/m1v1_nn.parquet'

# save the model to an in-memory buffer
buffer = BytesIO()
joblib.dump(m1v1_nn, buffer)

# upload the buffer to S3
buffer.seek(0)
s3.upload_fileobj(buffer, def_bucket, m1v1_nn_key_name)

```

```
# load the saved model from S3
#buffer = BytesIO()
#s3.download_fileobj(def_bucket, m1v1_nn_key_name, buffer)
#buffer.seek(0)
#m1v1_nn_fitted = joblib.load(buffer)
```

2.6.2 Lasso - Using GridSearchCV

```
[11]: # Start timer script
start_time = dt.datetime.today()

# Citation: Hochberg, 2018; Shanmukh, 2021
m2v1_ls_pip = Pipeline([('si', SimpleImputer(strategy='median')),
                        ('ss', StandardScaler()),
                        ('ls', Lasso(random_state=1699))])

alpha_hparam = [.01, .05, .1, .5, 1, 2]
selection_hparam = ['cyclic', 'random']

m2v1_ls_grd = {'ls__alpha': alpha_hparam,
              'ls__selection': selection_hparam
              }

m2v1_ls = GridSearchCV(m2v1_ls_pip,
                      m2v1_ls_grd,
                      scoring='neg_root_mean_squared_error',
                      n_jobs=2,
                      refit=True,
                      verbose=2)

m2v1_ls.fit(train_x01, train_y01)

print(f'Best Estimator:\n{m2v1_ls.best_estimator_}')
print(f'Coefficients:\n{m2v1_ls.best_estimator_.named_steps["ls"].coef_}')

print(pd.DataFrame(m2v1_ls.cv_results_))

train_m2v1_ls_y01_pred = m2v1_ls.predict(train_x01)
print(train_m2v1_ls_y01_pred)

test_m2v1_ls_y01_pred = m2v1_ls.predict(test_x01)
print(test_m2v1_ls_y01_pred)

# Display evaluation metrics
# R-sq
train_m2v1_ls_r2 = r2_score(train_y01, train_m2v1_ls_y01_pred)
```

```

test_m2v1_ls_r2 = r2_score(test_y01, test_m2v1_ls_y01_pred)

print(f'Train R-sq:\n{train_m2v1_ls_r2}')
print(f'Test R-sq:\n{test_m2v1_ls_r2}')

# RMSE
train_m2v1_ls_rmse = mean_squared_error(train_y01, train_m2v1_ls_y01_pred,
    ↪squared=False)
test_m2v1_ls_rmse = mean_squared_error(test_y01, test_m2v1_ls_y01_pred,
    ↪squared=False)

print(f'Train RMSE:\n{train_m2v1_ls_rmse}')
print(f'Test RMSE:\n{test_m2v1_ls_rmse}')

# End timer script
end_time = dt.datetime.today()
time_elapse = end_time - start_time
print(f'End Time = {end_time}')
print(f'Script Time = {time_elapse}')

```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 37 tasks      | elapsed: 15.4s
[Parallel(n_jobs=2)]: Done 60 out of 60 | elapsed: 18.6s finished

```

Best Estimator:

```

Pipeline(memory=None,
      steps=[('si',
              SimpleImputer(add_indicator=False, copy=True, fill_value=None,
                             missing_values=nan, strategy='median',
                             verbose=0)),
             ('ss',
              StandardScaler(copy=True, with_mean=True, with_std=True)),
             ('ls',
              Lasso(alpha=0.01, copy_X=True, fit_intercept=True,
                    max_iter=1000, normalize=False, positive=False,
                    precompute=False, random_state=1699, selection='random',
                    tol=0.0001, warm_start=False))],
      verbose=False)

```

Coefficients:

```

[-0.04698528  1.1655262  0.          -1.02395897  1.13308637 -0.
 -0.          0.          0.02814286 -0.          -0.0307876  0.
 0.          -0.          -0.01445511  0.          -0.          2.92859099
 2.50110937  5.60822761  1.40168756 -2.34296164 -1.18943262 -0.33822063
 -1.3657179  -3.4737577  -6.79353668  0.12912462  0.20841108 -0.04818647
 -0.51860673  2.15041078  0.72946348 -0.44590033  0.          -1.47779489
 -0.48312123  0.          1.54018429  0.49113008  0.56200594 -0.72542784

```

	-8.16308228	0.40673301	-1.2849073	0.46802939	0.73592459	1.6206294]
	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\	
0	1.025094	0.110517	0.009961	0.001149		
1	1.579347	0.095895	0.010434	0.000936		
2	0.512077	0.018247	0.010032	0.001147		
3	0.840451	0.166026	0.009896	0.001420		
4	0.431524	0.015982	0.008961	0.001188		
5	1.213542	0.502992	0.010119	0.001508		
6	0.287477	0.024440	0.008343	0.000368		
7	0.303767	0.005579	0.010572	0.001484		
8	0.253030	0.003281	0.010716	0.000799		
9	0.258238	0.008042	0.010476	0.000802		
10	0.242289	0.005520	0.010471	0.000613		
11	0.238470	0.028518	0.008936	0.001677		

	param_ls__alpha	param_ls__selection	\
0	0.01	cyclic	
1	0.01	random	
2	0.05	cyclic	
3	0.05	random	
4	0.1	cyclic	
5	0.1	random	
6	0.5	cyclic	
7	0.5	random	
8	1	cyclic	
9	1	random	
10	2	cyclic	
11	2	random	

	params	split0_test_score	\
0	{'ls__alpha': 0.01, 'ls__selection': 'cyclic'}	-11.378789	
1	{'ls__alpha': 0.01, 'ls__selection': 'random'}	-11.378654	
2	{'ls__alpha': 0.05, 'ls__selection': 'cyclic'}	-11.381059	
3	{'ls__alpha': 0.05, 'ls__selection': 'random'}	-11.380958	
4	{'ls__alpha': 0.1, 'ls__selection': 'cyclic'}	-11.396508	
5	{'ls__alpha': 0.1, 'ls__selection': 'random'}	-11.396574	
6	{'ls__alpha': 0.5, 'ls__selection': 'cyclic'}	-11.806394	
7	{'ls__alpha': 0.5, 'ls__selection': 'random'}	-11.806288	
8	{'ls__alpha': 1, 'ls__selection': 'cyclic'}	-12.008540	
9	{'ls__alpha': 1, 'ls__selection': 'random'}	-12.008700	
10	{'ls__alpha': 2, 'ls__selection': 'cyclic'}	-12.496541	
11	{'ls__alpha': 2, 'ls__selection': 'random'}	-12.496534	

	split1_test_score	split2_test_score	split3_test_score	\
0	-11.391382	-11.411164	-11.505775	
1	-11.391378	-11.411135	-11.505295	
2	-11.384965	-11.423042	-11.519926	
3	-11.384860	-11.423228	-11.520242	

4	-11.408320	-11.450277	-11.549378
5	-11.408500	-11.450586	-11.550226
6	-11.866156	-11.935607	-12.014084
7	-11.866341	-11.935859	-12.014298
8	-12.112724	-12.186260	-12.226056
9	-12.112880	-12.186494	-12.226203
10	-12.607476	-12.682111	-12.710059
11	-12.607472	-12.682084	-12.710063

	split4_test_score	mean_test_score	std_test_score	rank_test_score
0	-11.386550	-11.414732	0.046761	2
1	-11.386614	-11.414615	0.046588	1
2	-11.395211	-11.420840	0.051671	3
3	-11.395450	-11.420948	0.051800	4
4	-11.417847	-11.444466	0.055417	5
5	-11.418167	-11.444811	0.055679	6
6	-11.850882	-11.894625	0.072744	7
7	-11.851056	-11.894768	0.072833	8
8	-12.081994	-12.123115	0.076825	9
9	-12.082092	-12.123274	0.076840	10
10	-12.596496	-12.618537	0.074712	11
11	-12.596565	-12.618544	0.074706	12

```
[ 1.00885432  2.1155105 20.65181841 ...  5.04395528 20.70033313
  9.2996565 ]
[43.29645545  3.29328038 36.54323059 ... -0.46893988  1.17345827
  3.2687973 ]
```

Train R-sq:

0.6369229292317483

Test R-sq:

0.6479556475979478

Train RMSE:

11.389361452679568

Test RMSE:

11.339147217962033

End Time = 2023-04-13 19:29:57.540560

Script Time = 0:00:19.711830

```
[12]: coef_intercept = np.hstack((m2v1_ls.best_estimator_.named_steps["ls"].coef_,
                                m2v1_ls.best_estimator_.named_steps["ls"].
                                ↪intercept_))
# print(coef_intercept)

coef_intercept_df01 = pd.DataFrame(coef_intercept)
# display(coef_intercept_df01)

train_x01_col_names = list(train_x01.columns)
train_x01_col_names.append('intercept')
```

```

train_x01_col_names_df01 = pd.DataFrame(train_x01_col_names)
#display(train_x01_col_names_df01)

model_params = pd.concat([train_x01_col_names_df01, coef_intercept_df01],
↪axis=1)
display(model_params)

```

		0	0
0	borough_bronx	-0.046985	
1	borough_brooklyn	1.165526	
2	borough_manhattan	0.000000	
3	borough_queens	-1.023959	
4	borough_staten island	1.133086	
5	relative_data_year_-4	-0.000000	
6	relative_data_year_-3	-0.000000	
7	relative_data_year_-2	0.000000	
8	relative_data_year_-1	0.028143	
9	relative_data_year_0	-0.000000	
10	complaint_type_FELONY	-0.030788	
11	complaint_type_MISDEMEANOR	0.000000	
12	complaint_type_VIOLATION	0.000000	
13	annual_evictions_x_borough	-0.000000	
14	annual_complaint_counts	-0.014455	
15	annual_grad_n	0.000000	
16	annual_dropped_out_n	-0.000000	
17	totalpop	2.928591	
18	men	2.501109	
19	women	5.608228	
20	hispanic	1.401688	
21	white	-2.342962	
22	black	-1.189433	
23	native	-0.338221	
24	asian	-1.365718	
25	citizen	-3.473758	
26	income	-6.793537	
27	incomeerr	0.129125	
28	incomepercap	0.208411	
29	incomepercaperr	-0.048186	
30	professional	-0.518607	
31	service	2.150411	
32	office	0.729463	
33	construction	-0.445900	
34	production	0.000000	
35	drive	-1.477795	
36	carpool	-0.483121	
37	transit	0.000000	
38	walk	1.540184	

```

39             othertransp    0.491130
40             workathome     0.562006
41             meancommute    -0.725428
42             employed       -8.163082
43             privatework    0.406733
44             publicwork     -1.284907
45             selfemployed    0.468029
46             familywork     0.735925
47             unemployment    1.620629
48             intercept      24.475273

```

```

[13]: # specify the S3 bucket and key where you want to save the model
m2v1_ls_key_name = 'team_8_data/models/m2v1_ls.parquet'

# save the model to an in-memory buffer
buffer = BytesIO()
joblib.dump(m2v1_ls, buffer)

# upload the buffer to S3
buffer.seek(0)
s3.upload_fileobj(buffer, def_bucket, m2v1_ls_key_name)

# load the saved model from S3
#buffer = BytesIO()
#s3.download_fileobj(def_bucket, m2v1_ls_key_name, buffer)
#buffer.seek(0)
#m2v1_ls_fitted = joblib.load(buffer)

```

2.6.3 Lasso - Using BayesSearchCV

```

[14]: # Start timer script
start_time = dt.datetime.today()

# Citation: Hochberg, 2018; Shanmukh, 2021
m2v2_ls_pip = Pipeline([('si', SimpleImputer(strategy='median')),
                        ('ss', StandardScaler()),
                        ('ls', Lasso(random_state=1699))])

alpha_hparam = Real(1e-3, 1e3, prior='log-uniform')
selection_hparam = Categorical(['cyclic', 'random'])
max_iter_hparam = Integer(100, 5000, prior='log-uniform')
warm_start_hparam = Categorical([False, True])

m2v2_ls_grd = {'ls__alpha': alpha_hparam,
               'ls__selection': selection_hparam,
               'ls__max_iter': max_iter_hparam,

```

```

        'ls__warm_start': warm_start_hparam
    }

m2v2_ls = BayesSearchCV(m2v2_ls_pip,
                        m2v2_ls_grd,
                        scoring='neg_root_mean_squared_error',
                        cv=5,
                        n_jobs=2,
                        refit=True,
                        verbose=2)

m2v2_ls.fit(train_x01, train_y01)

print(f'Best Estimator:\n{m2v2_ls.best_estimator_}')
print(f'Coefficients:\n{m2v2_ls.best_estimator_.named_steps["ls"].coef_}')

print(pd.DataFrame(m2v2_ls.cv_results_))

train_m2v2_ls_y01_pred = m2v2_ls.predict(train_x01)
print(train_m2v2_ls_y01_pred)

test_m2v2_ls_y01_pred = m2v2_ls.predict(test_x01)
print(test_m2v2_ls_y01_pred)

# Display evaluation metrics
# R-sq
train_m2v2_ls_r2 = r2_score(train_y01, train_m2v2_ls_y01_pred)
test_m2v2_ls_r2 = r2_score(test_y01, test_m2v2_ls_y01_pred)

print(f'Train R-sq:\n{train_m2v2_ls_r2}')
print(f'Test R-sq:\n{test_m2v2_ls_r2}')

# RMSE
train_m2v2_ls_rmse = mean_squared_error(train_y01, train_m2v2_ls_y01_pred,
    ↪squared=False)
test_m2v2_ls_rmse = mean_squared_error(test_y01, test_m2v2_ls_y01_pred,
    ↪squared=False)

print(f'Train RMSE:\n{train_m2v2_ls_rmse}')
print(f'Test RMSE:\n{test_m2v2_ls_rmse}')

# End timer script
end_time = dt.datetime.today()
time_elapse = end_time - start_time
print(f'End Time = {end_time}')
print(f'Script Time = {time_elapse}')

```

```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s finished
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.6s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.6s finished
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s finished
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.8s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.8s finished
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.3s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.3s finished
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   3.4s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   3.4s finished
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.4s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.4s finished
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s finished
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```

```

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   0.7s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.6s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.6s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.7s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.3s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.3s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:  14.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:  14.7s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   2.9s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   2.9s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   7.5s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   7.5s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

```

```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.3s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.3s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   4.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   4.7s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   8.4s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   8.4s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   2.3s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   2.3s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   6.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   6.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:  26.6s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:  26.6s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```

```

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:  10.9s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:  10.9s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   2.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   2.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   5.7s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   5.7s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:  16.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:  16.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   4.2s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   4.2s finished

```



```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   3.5s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   3.5s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   8.4s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   8.4s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.4s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.4s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   2.9s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   2.9s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.0s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   1.1s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   3.8s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   3.8s finished

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```

```
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   14.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:   14.1s finished
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:    2.0s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:    2.0s finished
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:    1.1s remaining:   0.0s
[Parallel(n_jobs=2)]: Done   5 out of   5 | elapsed:    1.1s finished
```

Best Estimator:

```
Pipeline(memory=None,
          steps=[('si',
                  SimpleImputer(add_indicator=False, copy=True, fill_value=None,
                                missing_values=nan, strategy='median',
                                verbose=0)),
                 ('ss',
                  StandardScaler(copy=True, with_mean=True, with_std=True)),
                 ('ls',
                  Lasso(alpha=0.00312004499292689, copy_X=True,
                        fit_intercept=True, max_iter=5000, normalize=False,
                        positive=False, precompute=False, random_state=1699,
                        selection='cyclic', tol=0.0001, warm_start=False))],
          verbose=False)
```

Coefficients:

```
[-5.90772722e-02  1.20356958e+00  0.00000000e+00 -1.03729787e+00
  1.13319125e+00 -0.00000000e+00 -0.00000000e+00  2.94787994e-03
  3.38749310e-02 -2.05060234e-04 -2.82296287e-02  2.06826899e-02
  0.00000000e+00 -0.00000000e+00 -4.34594017e-02  0.00000000e+00
 -0.00000000e+00  8.48098006e+00  0.00000000e+00  2.76832909e+00
  6.97559873e-01 -3.30318362e+00 -2.14827171e+00 -3.71075424e-01
 -1.90796058e+00 -3.65510811e+00 -6.88025498e+00  1.54667377e-01
  4.43630919e-01 -1.76161247e-01 -1.81767382e+00  1.33115461e+00
  2.87452875e-01 -7.98665794e-01 -3.86055676e-01 -1.45597342e+00
 -4.86681209e-01  0.00000000e+00  1.52719040e+00  4.85589992e-01
  5.68145407e-01 -7.49278734e-01 -8.29828524e+00  0.00000000e+00
 -1.69877386e+00  2.42544215e-01  7.16688567e-01  1.60868297e+00]
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\
0	0.252393	0.044313	0.008819	0.001531	
1	0.211456	0.029397	0.009242	0.001724	
2	0.238254	0.037010	0.010291	0.002250	
3	0.241629	0.017929	0.010836	0.002084	
4	0.371777	0.079431	0.015980	0.004156	
5	1.241521	0.785646	0.011383	0.002698	

6	0.476834	0.089144	0.009877	0.001961
7	0.227144	0.035521	0.009860	0.002405
8	0.213959	0.030165	0.009917	0.002330
9	0.243137	0.040122	0.009101	0.001723
10	0.543308	0.085319	0.009564	0.001922
11	0.606364	0.112065	0.008954	0.001687
12	0.449924	0.079588	0.010595	0.003273
13	5.474890	2.152558	0.010317	0.002130
14	1.031375	0.176444	0.009314	0.001718
15	2.952774	1.177147	0.010025	0.002432
16	0.334927	0.052971	0.009441	0.001776
17	0.371395	0.064502	0.009407	0.001627
18	0.370219	0.069268	0.008443	0.001334
19	0.356148	0.062020	0.009524	0.002007
20	0.456968	0.099663	0.013050	0.004496
21	1.705872	0.324273	0.010758	0.004344
22	0.355035	0.064276	0.008949	0.001570
23	3.311975	0.950526	0.010320	0.002188
24	0.873525	0.198976	0.010325	0.003037
25	2.112508	0.328161	0.008992	0.002086
26	9.484036	1.248022	0.010280	0.001913
27	0.354692	0.058464	0.009997	0.001848
28	0.367649	0.061606	0.009237	0.001774
29	0.378164	0.070533	0.009635	0.002129
30	4.151578	2.864317	0.010427	0.002207
31	0.683975	0.111741	0.008902	0.001899
32	0.352353	0.063039	0.009234	0.001608
33	0.376071	0.064972	0.009555	0.001731
34	0.367897	0.065547	0.009819	0.002583
35	2.044964	0.413021	0.009135	0.002086
36	6.089054	1.929319	0.010797	0.002346
37	1.577792	0.276699	0.010048	0.001905
38	1.364861	0.264036	0.010090	0.001880
39	3.132879	0.511682	0.009999	0.002032
40	0.367394	0.062305	0.008894	0.001489
41	0.473200	0.077456	0.010350	0.002553
42	0.362873	0.056866	0.009705	0.002118
43	1.114455	0.251687	0.010146	0.002999
44	0.342149	0.052996	0.009398	0.001434
45	0.369063	0.061133	0.009580	0.001755
46	1.316128	0.174785	0.009341	0.001797
47	4.643917	1.957186	0.010486	0.002173
48	0.740760	0.162634	0.009555	0.002145
49	0.371850	0.062724	0.009669	0.001504

	param_ls__alpha	param_ls__max_iter	param_ls__selection	\
0	0.919229	372	cyclic	
1	20.785016	193	cyclic	

2	1.469648	441	random
3	9.380698	174	random
4	520.400979	2697	cyclic
5	0.002528	165	random
6	0.252225	949	cyclic
7	2.444393	4088	cyclic
8	107.97026	357	cyclic
9	0.859011	467	cyclic
10	0.020092	5000	cyclic
11	0.003409	246	random
12	0.056266	5000	cyclic
13	0.001	5000	cyclic
14	0.008623	5000	cyclic
15	0.001422	5000	cyclic
16	0.10612	100	random
17	0.001179	100	random
18	0.005593	100	cyclic
19	0.030282	100	random
20	0.012457	100	random
21	0.006117	5000	random
22	0.040493	100	cyclic
23	0.00312	5000	cyclic
24	0.040469	5000	random
25	0.004728	5000	cyclic
26	0.001	5000	random
27	0.061997	100	random
28	0.016873	100	random
29	0.001819	100	random
30	0.00183	5000	cyclic
31	0.013661	5000	cyclic
32	0.046368	100	random
33	0.006918	100	random
34	0.001	100	cyclic
35	0.004217	5000	random
36	0.001	5000	cyclic
37	0.010632	5000	random
38	0.015439	5000	random
39	0.002274	5000	random
40	0.002292	100	cyclic
41	0.033243	5000	cyclic
42	0.009614	100	cyclic
43	0.023583	5000	random
44	0.080033	100	cyclic
45	0.026002	100	cyclic
46	0.007288	5000	cyclic
47	0.001158	5000	cyclic
48	0.080142	5000	random
49	0.003943	100	cyclic

	param_ls__warm_start	params \
0	True	{'ls__alpha': 0.9192293938041608, 'ls__max_ite...
1	True	{'ls__alpha': 20.785016353070272, 'ls__max_ite...
2	True	{'ls__alpha': 1.4696478744138841, 'ls__max_ite...
3	True	{'ls__alpha': 9.380698318075792, 'ls__max_iter...
4	True	{'ls__alpha': 520.4009792978661, 'ls__max_iter...
5	True	{'ls__alpha': 0.0025276741815852596, 'ls__max_...
6	True	{'ls__alpha': 0.2522250210775078, 'ls__max_ite...
7	True	{'ls__alpha': 2.4443925994929305, 'ls__max_ite...
8	False	{'ls__alpha': 107.97025991821005, 'ls__max_ite...
9	False	{'ls__alpha': 0.8590108805849646, 'ls__max_ite...
10	False	{'ls__alpha': 0.020091744449482356, 'ls__max_i...
11	True	{'ls__alpha': 0.003409256952063032, 'ls__max_i...
12	True	{'ls__alpha': 0.056266004320859396, 'ls__max_i...
13	False	{'ls__alpha': 0.001, 'ls__max_iter': 5000, 'ls...
14	True	{'ls__alpha': 0.008623045081308458, 'ls__max_i...
15	False	{'ls__alpha': 0.0014221967527329972, 'ls__max_...
16	False	{'ls__alpha': 0.10611954673637193, 'ls__max_it...
17	True	{'ls__alpha': 0.0011791282733941873, 'ls__max_...
18	False	{'ls__alpha': 0.005592992209800713, 'ls__max_i...
19	True	{'ls__alpha': 0.030281735565660575, 'ls__max_i...
20	False	{'ls__alpha': 0.012457400581715239, 'ls__max_i...
21	False	{'ls__alpha': 0.006117246602685628, 'ls__max_i...
22	False	{'ls__alpha': 0.04049343717819463, 'ls__max_it...
23	False	{'ls__alpha': 0.00312004499292689, 'ls__max_it...
24	False	{'ls__alpha': 0.04046875004317316, 'ls__max_it...
25	True	{'ls__alpha': 0.004728162592110137, 'ls__max_i...
26	False	{'ls__alpha': 0.001, 'ls__max_iter': 5000, 'ls...
27	False	{'ls__alpha': 0.06199705881601676, 'ls__max_it...
28	True	{'ls__alpha': 0.01687341186753009, 'ls__max_it...
29	False	{'ls__alpha': 0.0018185560271089783, 'ls__max_...
30	True	{'ls__alpha': 0.0018303658598422532, 'ls__max_...
31	True	{'ls__alpha': 0.013661350637267831, 'ls__max_i...
32	True	{'ls__alpha': 0.04636797353331355, 'ls__max_it...
33	True	{'ls__alpha': 0.006918445469792199, 'ls__max_i...
34	False	{'ls__alpha': 0.001, 'ls__max_iter': 100, 'ls...
35	False	{'ls__alpha': 0.00421690504388654, 'ls__max_it...
36	True	{'ls__alpha': 0.001, 'ls__max_iter': 5000, 'ls...
37	True	{'ls__alpha': 0.010632072820022796, 'ls__max_i...
38	False	{'ls__alpha': 0.015438722672047024, 'ls__max_i...
39	False	{'ls__alpha': 0.0022740964287695036, 'ls__max_...
40	False	{'ls__alpha': 0.0022923216227139103, 'ls__max_...
41	True	{'ls__alpha': 0.03324342462416447, 'ls__max_it...
42	False	{'ls__alpha': 0.009614323672543546, 'ls__max_i...
43	True	{'ls__alpha': 0.02358308086767287, 'ls__max_it...
44	True	{'ls__alpha': 0.08003292102153944, 'ls__max_it...
45	False	{'ls__alpha': 0.02600213461730265, 'ls__max_it...

```

46         False {'ls__alpha': 0.0072884291262491525, 'ls__max_...
47         False {'ls__alpha': 0.0011575416936216076, 'ls__max_...
48         True  {'ls__alpha': 0.08014202444546731, 'ls__max_it...
49         False {'ls__alpha': 0.003942785721175991, 'ls__max_i...

```

	split0_test_score	split1_test_score	split2_test_score \
0	-11.977457	-12.082645	-12.155722
1	-18.828273	-18.816528	-18.873667
2	-12.217014	-12.327770	-12.406181
3	-16.469237	-16.476761	-16.486357
4	-18.828273	-18.816528	-18.873667
5	-11.380662	-11.395944	-11.410907
6	-11.507037	-11.556216	-11.604021
7	-12.747390	-12.846645	-12.912869
8	-18.828273	-18.816528	-18.873667
9	-11.955538	-12.058478	-12.134445
10	-11.380589	-11.386755	-11.415462
11	-11.379845	-11.396099	-11.409756
12	-11.381898	-11.386847	-11.425501
13	-11.379818	-11.398849	-11.406317
14	-11.378792	-11.392354	-11.410356
15	-11.379706	-11.398501	-11.406496
16	-11.400669	-11.413166	-11.456496
17	-11.382560	-11.395759	-11.412853
18	-11.380164	-11.395063	-11.410263
19	-11.380880	-11.382676	-11.418229
20	-11.381084	-11.389948	-11.414542
21	-11.378914	-11.394251	-11.408826
22	-11.381145	-11.383385	-11.418390
23	-11.379440	-11.396892	-11.407253
24	-11.380317	-11.382732	-11.420022
25	-11.379224	-11.395523	-11.408173
26	-11.379918	-11.398930	-11.406477
27	-11.384259	-11.389523	-11.429869
28	-11.381018	-11.388484	-11.415529
29	-11.382439	-11.395379	-11.413013
30	-11.379603	-11.398171	-11.406777
31	-11.379134	-11.389215	-11.413041
32	-11.381865	-11.384254	-11.423136
33	-11.381857	-11.392092	-11.413386
34	-11.379834	-11.397019	-11.409298
35	-11.379235	-11.395867	-11.407826
36	-11.379818	-11.398849	-11.406317
37	-11.378559	-11.390977	-11.411400
38	-11.379298	-11.388353	-11.414127
39	-11.379691	-11.397716	-11.407038
40	-11.379802	-11.396319	-11.409414
41	-11.380290	-11.382304	-11.417945

42	-11.380949	-11.394112	-11.411515
43	-11.380230	-11.385291	-11.416149
44	-11.388175	-11.396595	-11.436773
45	-11.381719	-11.385534	-11.414868
46	-11.378868	-11.393370	-11.409555
47	-11.379802	-11.398752	-11.406384
48	-11.388123	-11.396783	-11.437692
49	-11.379968	-11.395684	-11.409814

	split3_test_score	split4_test_score	mean_test_score	std_test_score \
0	-12.199795	-12.052402	-12.093604	0.078029
1	-18.934900	-19.055553	-18.901784	0.087412
2	-12.424359	-12.291528	-12.333370	0.076028
3	-16.600301	-16.742277	-16.554987	0.105179
4	-18.934900	-19.055553	-18.901784	0.087412
5	-11.504640	-11.387121	-11.415855	0.045537
6	-11.705059	-11.555625	-11.585592	0.067147
7	-12.945755	-12.854852	-12.861502	0.067858
8	-18.934900	-19.055553	-18.901784	0.087412
9	-12.179510	-12.029071	-12.071409	0.078797
10	-11.510673	-11.389846	-11.416665	0.048486
11	-11.503854	-11.386609	-11.415232	0.045434
12	-11.522605	-11.397110	-11.422792	0.052142
13	-11.500890	-11.385568	-11.414288	0.044305
14	-11.504997	-11.386338	-11.414567	0.046402
15	-11.501125	-11.385462	-11.414258	0.044446
16	-11.555529	-11.423823	-11.449937	0.055954
17	-11.506265	-11.388701	-11.417228	0.045658
18	-11.503716	-11.386209	-11.415083	0.045460
19	-11.513630	-11.391245	-11.417332	0.049969
20	-11.509086	-11.388563	-11.416645	0.047573
21	-11.503610	-11.385897	-11.414300	0.045754
22	-11.515097	-11.391125	-11.417828	0.050411
23	-11.502023	-11.385352	-11.414192	0.044947
24	-11.516632	-11.393011	-11.418543	0.051028
25	-11.502993	-11.385667	-11.414316	0.045402
26	-11.501023	-11.385705	-11.414411	0.044312
27	-11.526411	-11.399865	-11.425986	0.052637
28	-11.510211	-11.389163	-11.416881	0.048111
29	-11.506586	-11.388744	-11.417232	0.045829
30	-11.501183	-11.385397	-11.414226	0.044508
31	-11.507987	-11.387405	-11.415356	0.047671
32	-11.519297	-11.394784	-11.420667	0.051450
33	-11.508549	-11.388722	-11.416921	0.047009
34	-11.503119	-11.385764	-11.415007	0.045193
35	-11.502647	-11.385574	-11.414230	0.045258
36	-11.500890	-11.385568	-11.414288	0.044305
37	-11.505984	-11.386739	-11.414732	0.046892

38	-11.508911	-11.388004	-11.415739	0.048022
39	-11.501598	-11.385604	-11.414330	0.044654
40	-11.503101	-11.385844	-11.414896	0.045235
41	-11.514004	-11.391430	-11.417194	0.050231
42	-11.504158	-11.387072	-11.415561	0.045465
43	-11.511637	-11.390373	-11.416736	0.049037
44	-11.533929	-11.406318	-11.432358	0.053377
45	-11.511093	-11.389242	-11.416491	0.048708
46	-11.504270	-11.386175	-11.414448	0.046047
47	-11.500968	-11.385536	-11.414288	0.044347
48	-11.536468	-11.407262	-11.433265	0.054250
49	-11.503447	-11.385891	-11.414961	0.045382

	rank_test_score
0	44
1	48
2	45
3	47
4	48
5	23
6	42
7	46
8	48
9	43
10	26
11	19
12	37
13	5
14	13
15	4
16	41
17	31
18	18
19	33
20	25
21	8
22	34
23	1
24	35
25	9
26	11
27	38
28	28
29	32
30	2
31	20
32	36
33	29


```

34          17
35          3
36          5
37         14
38         22
39         10
40         15
41         30
42         21
43         27
44         39
45         24
46         12
47          7
48         40
49         16
[ 0.80813214  2.15188504 20.47191416 ...  4.87480063 20.68240508
 9.20470403]
[43.3663052  3.55764009 36.54084305 ... -0.57203178  1.09980524
 3.14855382]
Train R-sq:
0.6370404443324977
Test R-sq:
0.6480368310543366
Train RMSE:
11.387518138430686
Test RMSE:
11.337839706231485
End Time = 2023-04-13 19:34:32.039668
Script Time = 0:04:34.328540

```

```

[15]: coef_intercept = np.hstack((m2v2_ls.best_estimator_.named_steps["ls"].coef_,
                                   m2v2_ls.best_estimator_.named_steps["ls"].
                                   ↪intercept_))
      #print(coef_intercept)

      coef_intercept_df01 = pd.DataFrame(coef_intercept)
      #display(coef_intercept_df01)

      train_x01_col_names = list(train_x01.columns)
      train_x01_col_names.append('intercept')

      train_x01_col_names_df01 = pd.DataFrame(train_x01_col_names)
      #display(train_x01_col_names_df01)

      model_params = pd.concat([train_x01_col_names_df01, coef_intercept_df01],
                                ↪axis=1)

```

```
display(model_params)
```

	0	0
0	borough_bronx	-0.059077
1	borough_brooklyn	1.203570
2	borough_manhattan	0.000000
3	borough_queens	-1.037298
4	borough_staten_island	1.133191
5	relative_data_year_-4	-0.000000
6	relative_data_year_-3	-0.000000
7	relative_data_year_-2	0.002948
8	relative_data_year_-1	0.033875
9	relative_data_year_0	-0.000205
10	complaint_type_FELONY	-0.028230
11	complaint_type_MISDEMEANOR	0.020683
12	complaint_type_VIOLATION	0.000000
13	annual_evictions_x_borough	-0.000000
14	annual_complaint_counts	-0.043459
15	annual_grad_n	0.000000
16	annual_dropped_out_n	-0.000000
17	totalpop	8.480980
18	men	0.000000
19	women	2.768329
20	hispanic	0.697560
21	white	-3.303184
22	black	-2.148272
23	native	-0.371075
24	asian	-1.907961
25	citizen	-3.655108
26	income	-6.880255
27	incomeerr	0.154667
28	incomepercap	0.443631
29	incomepercaperr	-0.176161
30	professional	-1.817674
31	service	1.331155
32	office	0.287453
33	construction	-0.798666
34	production	-0.386056
35	drive	-1.455973
36	carpool	-0.486681
37	transit	0.000000
38	walk	1.527190
39	othertransp	0.485590
40	workathome	0.568145
41	meancommute	-0.749279
42	employed	-8.298285
43	privatework	0.000000
44	publicwork	-1.698774

```

45             selfemployed    0.242544
46             familywork     0.716689
47             unemployment    1.608683
48             intercept      24.475273

```

```

[16]: # specify the S3 bucket and key where you want to save the model
m2v2_ls_key_name = 'team_8_data/models/m2v2_ls.parquet'

# save the model to an in-memory buffer
buffer = BytesIO()
joblib.dump(m2v2_ls, buffer)

# upload the buffer to S3
buffer.seek(0)
s3.upload_fileobj(buffer, def_bucket, m2v2_ls_key_name)

# load the saved model from S3
#buffer = BytesIO()
#s3.download_fileobj(def_bucket, m2v2_ls_key_name, buffer)
#buffer.seek(0)
#m2v2_ls_fitted = joblib.load(buffer)

```

2.7 Release Resources

```

[17]: %%html

<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:
↪shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
    els = document.getElementsByClassName("sm-command-button");
    els[0].click();
}
catch(err) {
    // NoOp
}
</script>

```

<IPython.core.display.HTML object>

```

[18]: %%javascript

try {
    Jupyter.notebook.save_checkpoint();
    Jupyter.notebook.session.delete();
}

```

```
catch(err) {  
    // NoOp  
}
```

<IPython.core.display.Javascript object>